

Report for HPC LAB

Name: Bhaskar R

Roll No: CED18I009

Programming Environment: MPI

Problem: Vector Multiplication

Date: 22nd October 2021

Hardware Configuration:

CPU NAME : Intel core i5 – 8250U @ 1.60 GHz

Number of Sockets : 1

Cores per Socket : 4

Threads per core : 2

L1 Cache size : 32KB (Per Core)

L2 Cache size : 256KB (Per Core)

L3 Cache size : 6MB (Shared)

RAM : 8 GB

Serial Code:

```
#include "mpi.h"
#include <stdio.h>
#include <stdlib.h>
#define n 100
#define m 100
int main()
{
    double a[n], b[n], c[n];
    float startTime, endTime, execTime;
    int i, k;
    float rtime;
    startTime = MPI_WTIME();
    for (i = 0; i < n; i++)
    {
        a[i] = i * 10.236;    // Use Random function and assign a[i]
        b[i] = i * 152.123; // Use Random function and assign b[i]
        for (int j = 0; j < m; j++)
            c[i] = a[i] * b[i];
        //printf("The value of a[%d] = %lf and b[%d] = %lf and result c[%d] = %lf done\n", i, a[i], i, b[i], i, c[i], omp_rank);
    }
    endTime = MPI_WTIME();
    execTime = endTime - startTime;
    rtime = execTime;
    printf("\n rtime=%f\n", rtime);
    return (0);
}
```

Parallel Code :

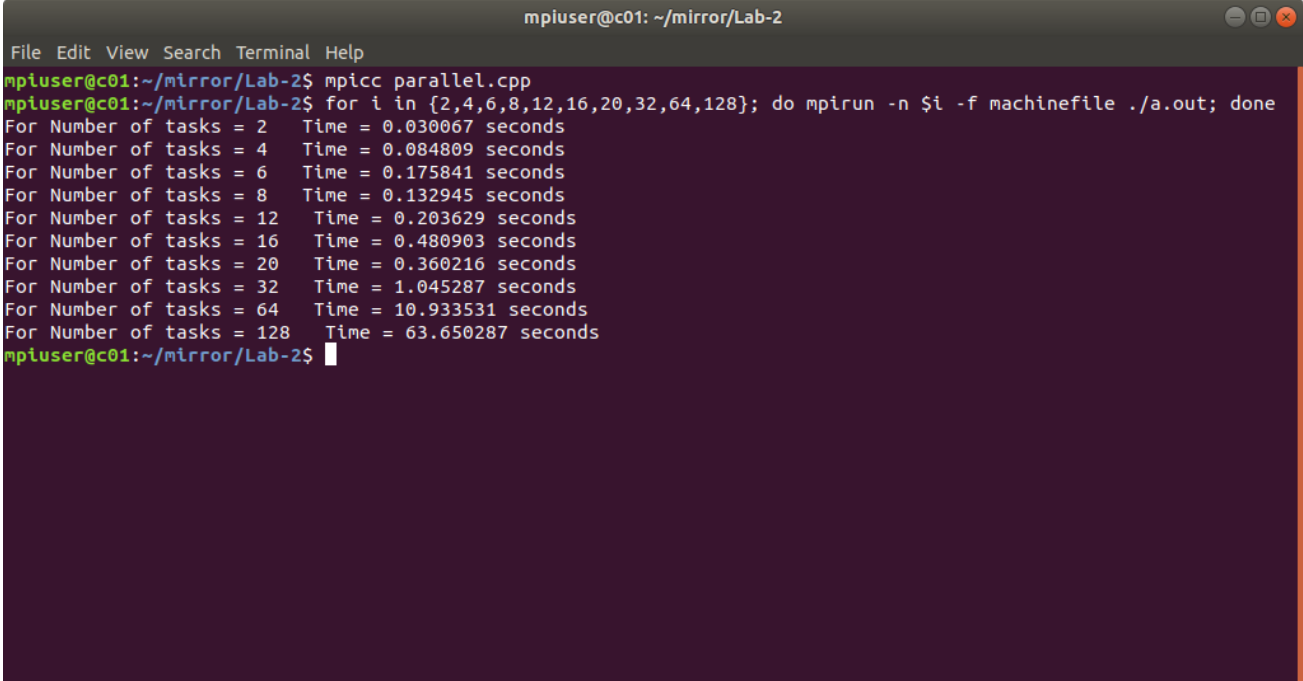
```
#include "mpi.h"
#include <stdio.h>
#include <stdlib.h>
#define array_size 100
#define MASTER 0
#define FROM_MASTER 1
#define FROM_WORKER 2
int main(int argc, char *argv[])
{
    int numtasks, taskid, numworkers, source, dest, mtype, segment, aveseq, extra,
        offset, i, j, k, rc;
    double starttime, endtime;
    long double a[array_size], b[array_size], c[array_size];
    MPI_Status status;
    MPI_Init(&argc, &argv);
    starttime = MPI_Wtime();
    MPI_Comm_rank(MPI_COMM_WORLD, &taskid);
    MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
    if (numtasks < 2)
    {
        printf("Need atleast two MPI tasks. Quitting...\n");
        MPI_Abort(MPI_COMM_WORLD, rc);
        exit(1);
    }
    numworkers = numtasks - 1;
    //master task:
    if (taskid == MASTER)
    {
        for (i = 0; i < array_size; i++)
            a[i] = i;
        for (j = 0; j < array_size; j++)
            b[j] = j;
        aveseq = array_size / numworkers;
        extra = array_size % numworkers;
        offset = 0;
        mtype = FROM_MASTER;
        for (dest = 1; dest <= numworkers; dest++)
        {
            segment = (dest <= extra) ? aveseq + 1 : aveseq;
            MPI_Send(&offset, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);
            MPI_Send(&segment, 1, MPI_INT, dest, mtype,
MPI_COMM_WORLD);
            MPI_Send(&a[offset], segment, MPI_LONG_DOUBLE, dest, mtype,
MPI_COMM_WORLD);
            MPI_Send(&b[offset], segment, MPI_LONG_DOUBLE, dest, mtype,
MPI_COMM_WORLD);
            offset = offset + segment;
        }
        //receive from worker:
        mtype = FROM_WORKER;
```

```

        for (i = 1; i <= numworkers; i++)
        {
            source = i;
            MPI_Recv(&offset, 1, MPI_INT, source, mtype, MPI_COMM_WORLD,
&status);
            MPI_Recv(&segment, 1, MPI_INT, source, mtype,
MPI_COMM_WORLD,
                        &status);
            MPI_Recv(&c[offset], segment, MPI_LONG_DOUBLE, source, mtype,
MPI_COMM_WORLD, &status);
        }
        printf("\nResultant Vector:\n");
        for (i = 0; i < array_size; i++)
            printf("%6.2Lf ", c[i]);
        endtime = MPI_Wtime();
        printf("That took %f seconds\n", endtime - starttime);
        printf("\nDone.\n");
    }
    //Worker task:
    if (taskid > MASTER)
    {
        mtype = FROM_MASTER;
        MPI_Recv(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD,
&status);
        MPI_Recv(&segment, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD,
&status);
        MPI_Recv(&a, segment, MPI_LONG_DOUBLE, MASTER, mtype,
MPI_COMM_WORLD, &status);
        MPI_Recv(&b, segment, MPI_LONG_DOUBLE, MASTER, mtype,
MPI_COMM_WORLD, &status);
        //mat addition
        for (i = 0; i < segment; i++)
            c[i] = a[i] * b[i];
        mtype = FROM_WORKER;
        MPI_Send(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);
        MPI_Send(&segment, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);
        MPI_Send(&c, segment, MPI_LONG_DOUBLE, MASTER, mtype,
MPI_COMM_WORLD);
    }
    MPI_Finalize();
}

```

Output :



```
mpiuser@c01: ~/mirror/Lab-2
File Edit View Search Terminal Help
mpiuser@c01:~/mirror/Lab-2$ mpicc parallel.cpp
mpiuser@c01:~/mirror/Lab-2$ for i in {2,4,6,8,12,16,20,32,64,128}; do mpirun -n $i -f machinefile ./a.out; done
For Number of tasks = 2    Time = 0.030067 seconds
For Number of tasks = 4    Time = 0.084809 seconds
For Number of tasks = 6    Time = 0.175841 seconds
For Number of tasks = 8    Time = 0.132945 seconds
For Number of tasks = 12   Time = 0.203629 seconds
For Number of tasks = 16   Time = 0.480903 seconds
For Number of tasks = 20   Time = 0.360216 seconds
For Number of tasks = 32   Time = 1.045287 seconds
For Number of tasks = 64   Time = 10.933531 seconds
For Number of tasks = 128  Time = 63.650287 seconds
mpiuser@c01:~/mirror/Lab-2$
```

Compilation and Execution:

Compiling using mpic++ parallel.cpp

For execution use

```
for i in {2,4,6,8,12,16,20,32,64,128}; do mpirun -n $i -f machinefile ./a.out
```

Observations:

Number of Threads	Execution Time	Speed-up	Parallelization Fraction
1	0.01	1	
2	0.03	0.3333	-4.0006
4	0.08	0.125	-9.3333
6	0.17	0.0588	-19.2082
8	0.13	0.0769	-13.7187
12	0.20	0.05	-20.7273
16	0.48	0.0208	-50.2154
20	0.36	0.0278	-36.8118
32	1.04	0.0096	-106.4946
64	10.9	0.0009	-1127.7319
128	63.65	0.0002	-5038.3622

Speed up can be found using the following formula,

$$S(n)=T(1)/T(n)$$

where, $S(n)$ = Speedup for thread count 'n'

$T(1)$ = Execution Time for Thread count '1' (serial code)

$T(n)$ = Execution Time for Thread count 'n' (serial code)


Parallelization Fraction can be found using the following formula, $S(n)=1/((1 - p) + p/n)$

where, $S(n)$ = Speedup for thread count 'n'


n = Number of threads

p = Parallelization fraction

Number of Threads vs Execution Time:



Number of Threads vs Speed Up:



Inference:

- Execution time increases with an increase in the number of threads

Since the problem is of smaller complexity the overheads of parallelization seem to have more effects here.