

Report for HPC LAB

Name: Bhaskar R

Roll No: CED18I009

Programming Environment: MPI

Problem: Matrix Addition

Date: 22nd October 2021

Hardware Configuration:

CPU NAME : Intel core i5 – 8250U @ 1.60 GHz

Number of Sockets : 1

Cores per Socket : 4

Threads per core : 2

L1 Cache size : 32KB (Per Core)

L2 Cache size : 256KB (Per Core)

L3 Cache size : 6MB (Shared)

RAM : 8 GB

Serial Code:

```
#include "mpi.h"
#include <stdio.h>
#include <stdlib.h>
#define NR 5
#define NC 5
#define MASTER 0
#define FROM_MASTER 1
#define FROM_WORKER 2
int main()
{
    double start, end;
    int i, j;
    long double a[NR][NC], b[NR][NC], c[NR][NC];
    start = MPI_Wtime();
    for (i = 0; i < NR; i++)
        for (j = 0; j < NC; j++)
            a[i][j] = i + j * 1.785;
    for (i = 0; i < NR; i++)
        for (j = 0; j < NC; j++)
            b[i][j] = i + j * 0.987;
    for (i = 0; i < NR; i++)
        for (j = 0; j < NC; j++)
        {
            c[i][j] = a[i][j] + b[i][j];
        }
    printf("\nResultant Matrix:\n");
    for (i = 0; i < NR; i++)
    {
        printf("\n");
        for (j = 0; j < NC; j++)
            printf("%3.1Lf ", c[i][j]);
    }
    printf("\nFinished.\n");
}
```

```

    end = MPI_Wtime();
    printf("\nTime= %f", end - start);
    return 0;
}

```

Parallel Code :

```

#include "mpi.h"
#include <stdio.h>
#include <stdlib.h>
#define NR 5
#define NC 5
#define MASTER 0
#define FROM_MASTER 1
#define FROM_WORKER 2
// cluster parallel code
int main(int argc, char *argv[])
{
    int numtasks, taskid, numworkers, source, dest, mtype, rows, averow, extra, offset, i,
    j,
        k, rc;
    long double a[NR][NC], b[NC][NC], c[NR][NC];
    MPI_Status status;
    double start, end;
    MPI_Init(&argc, &argv);
    start = MPI_Wtime();
    MPI_Comm_rank(MPI_COMM_WORLD, &taskid);
    MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
    if (numtasks < 2)
    {
        printf("Need atleast two MPI tasks. Quitting...\n");
        MPI_Abort(MPI_COMM_WORLD, rc);
        exit(1);
    }
    char pro_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(pro_name, &name_len);
    printf("-From from %s, rank %d, out of %d processors\n", pro_name, taskid,
        numtasks);
    numworkers = numtasks - 1;
    // master task:
    if (taskid == MASTER)
    {
        for (i = 0; i < NR; i++)
            for (j = 0; j < NC; j++)
                a[i][j] = i + j * 1.785;
        for (i = 0; i < NC; i++)
            for (j = 0; j < NC; j++)
                b[i][j] = i * j * 0.897;
        averow = NR / numworkers;
        extra = NR % numworkers;
        offset = 0;
    }
}

```

```

mtype = FROM_MASTER;
for (dest = 1; dest <= numworkers; dest++)
{
    rows = (dest <= extra) ? averow + 1 : averow;
    MPI_Send(&offset, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);
    MPI_Send(&rows, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);
    MPI_Send(&a[offset][0], rows * NC, MPI_LONG_DOUBLE, dest,
mtype,
        MPI_COMM_WORLD);
    MPI_Send(&b[offset][0], rows * NC, MPI_LONG_DOUBLE, dest,
mtype,
        MPI_COMM_WORLD);
    offset = offset + rows;
}
// receive from worker:
mtype = FROM_WORKER;
for (i = 1; i <= numworkers; i++)
{
    source = i;
    MPI_Recv(&offset, 1, MPI_INT, source, mtype, MPI_COMM_WORLD,
&status);
    MPI_Recv(&rows, 1, MPI_INT, source, mtype, MPI_COMM_WORLD,
&status);
    MPI_Recv(&c[offset][0], rows * NC, MPI_LONG_DOUBLE, source,
mtype, MPI_COMM_WORLD, &status);
}
printf("\nResultant Matrix:\n");
for (i = 0; i < NR; i++)
{
    printf("\n");
    for (j = 0; j < NC; j++)
        printf("%6.2Lf ", c[i][j]);
}
printf("\nDone.\n");
end = MPI_Wtime();
printf("\nTime= %f", end - start);
}
// Worker task:
if (taskid > MASTER)
{
    mtype = FROM_MASTER;
    MPI_Recv(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD,
&status);
    MPI_Recv(&rows, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD,
&status);
    MPI_Recv(&a, rows * NC, MPI_LONG_DOUBLE, MASTER, mtype,
        MPI_COMM_WORLD, &status);
    MPI_Recv(&b, rows * NC, MPI_LONG_DOUBLE, MASTER, mtype,
        MPI_COMM_WORLD, &status);
    // mat addition
    for (k = 0; k < NC; k++)
        for (i = 0; i < rows; i++)
        {

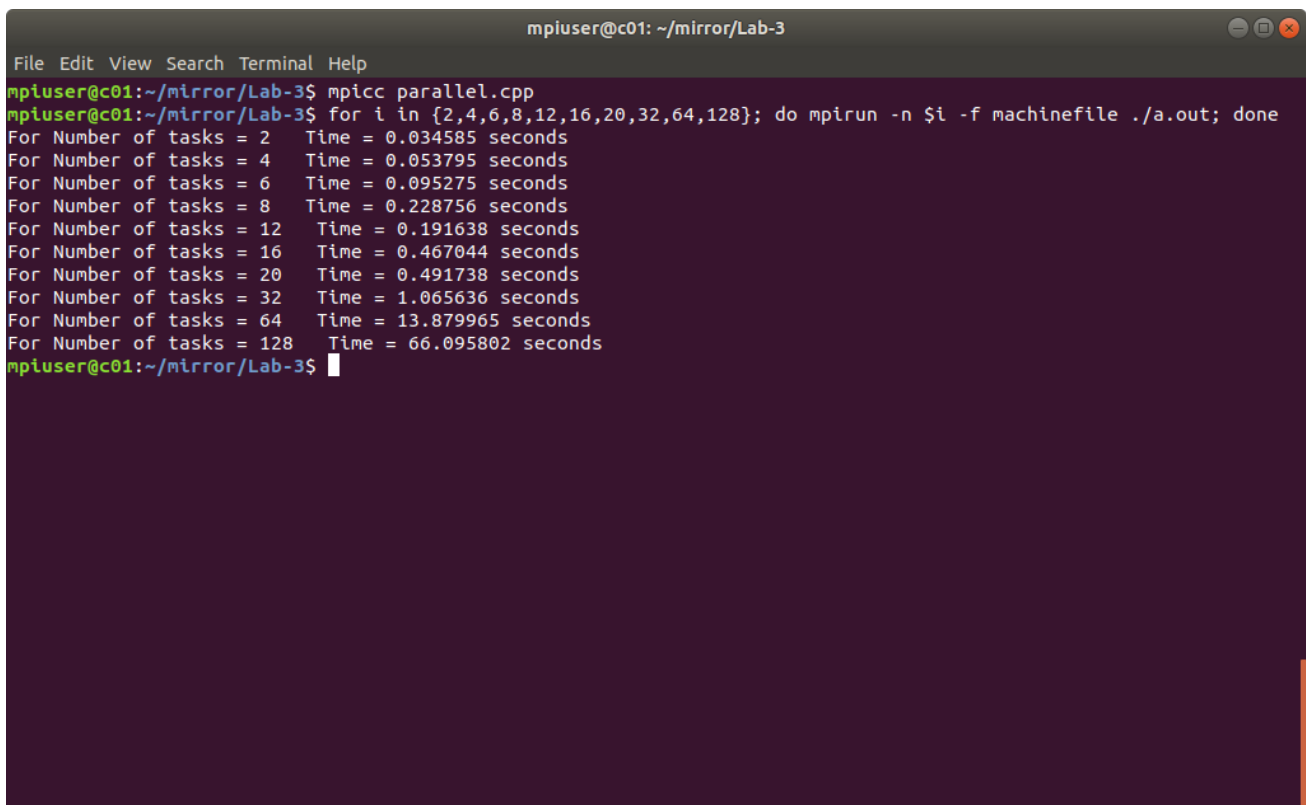
```

```

        c[i][k] = a[i][k] + b[i][k];
    }
    mtype = FROM_WORKER;
    MPI_Send(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);
    MPI_Send(&rows, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);
    MPI_Send(&c, rows * NC, MPI_LONG_DOUBLE, MASTER, mtype,
MPI_COMM_WORLD);
}
MPI_Finalize();
}

```

Output :



A terminal window titled 'mpiuser@c01: ~/mirror/Lab-3' displays the output of an MPI program. The user has compiled 'parallel.cpp' using 'mpicc' and then executed a loop of tests for task counts from 2 to 128. The output shows a linear increase in execution time as the number of tasks increases, with a significant jump at 64 tasks.

```

mpiuser@c01:~/mirror/Lab-3$ mpicc parallel.cpp
mpiuser@c01:~/mirror/Lab-3$ for i in {2,4,6,8,12,16,20,32,64,128}; do mpirun -n $i -f machinefile ./a.out; done
For Number of tasks = 2      Time = 0.034585 seconds
For Number of tasks = 4      Time = 0.053795 seconds
For Number of tasks = 6      Time = 0.095275 seconds
For Number of tasks = 8      Time = 0.228756 seconds
For Number of tasks = 12     Time = 0.191638 seconds
For Number of tasks = 16     Time = 0.467044 seconds
For Number of tasks = 20     Time = 0.491738 seconds
For Number of tasks = 32     Time = 1.065636 seconds
For Number of tasks = 64     Time = 13.879965 seconds
For Number of tasks = 128    Time = 66.095802 seconds
mpiuser@c01:~/mirror/Lab-3$

```

Compilation and Execution:

Compiling using mpic++ parallel.cpp

For execution use

```
for i in {2,4,6,8,12,16,20,32,64,128}; do mpirun -n $i -f machinefile ./a.out
```

Observations:

Number of Threads	Execution Time	Speed-up	Parallelization Fraction
1	0.01	1	
2	0.03	0.3333	-4.0006
4	0.05	0.2	-5.3333
6	0.09	0.1111	-9.6011
8	0.22	0.0455	-23.9749
12	0.19	0.0526	-19.6488
16	0.46	0.0217	-48.0885
20	0.49	0.0204	-50.547
32	1.06	0.0094	-108.7824
64	13.87	0.0007	-1450.2313
128	66.09	0.0002	-5038.3622

Speed up can be found using the following formula,

$$S(n)=T(1)/T(n)$$

where, S(n) = Speedup for thread count 'n'

T(1) = Execution Time for Thread count '1' (serial code)

T(n) = Execution Time for Thread count 'n' (serial code)

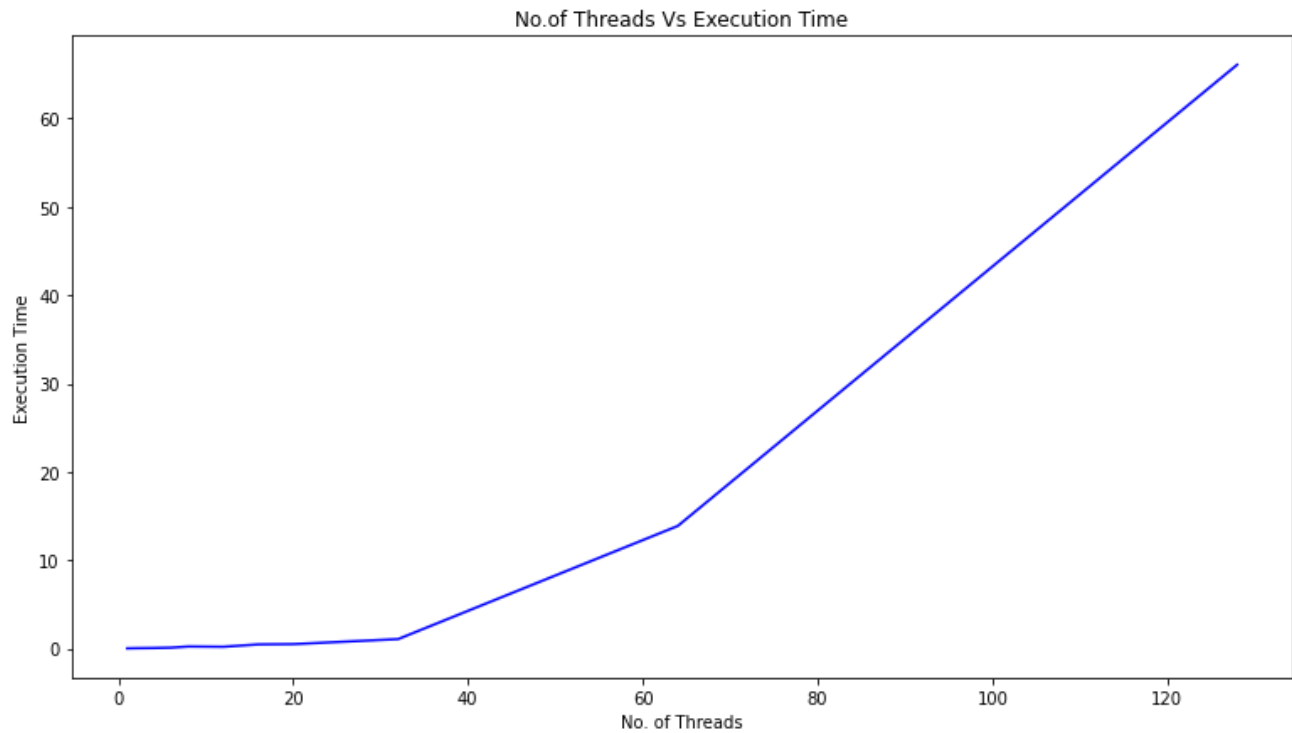
Parallelization Fraction can be found using the following formula, $S(n)=1/((1 - p) + p/n)$

where, S(n) = Speedup for thread count 'n'


n = Number of threads

p = Parallelization fraction

Number of Threads vs Execution Time:



Number of Threads vs Speed Up:



Inference:

- Execution time is increasing with an increase in the number of threads, But at 6 threads less execution time occurred. Since the problem is of smaller complexity the overheads of parallelization seem to have more effects here.