

# Report for HPC LAB

**Name:** Bhaskar R

**Roll No:** CED18I009

**Programming Environment:** MPI

**Problem:** Matrix Multiplication

**Date:** 22<sup>nd</sup> October 2021

## Hardware Configuration:

CPU NAME : Intel core i5 – 8250U @ 1.60 GHz

Number of Sockets : 1

Cores per Socket : 4

Threads per core : 2

L1 Cache size : 32KB (Per Core)

L2 Cache size : 256KB (Per Core)

L3 Cache size : 6MB (Shared)

RAM : 8 GB

## Serial Code:

```
#include "mpi.h"
#include <stdio.h>
#include <stdlib.h>
#define NR 5
#define NC 5
#define MASTER 0
#define FROM_MASTER 1
#define FROM_WORKER 2
int main()
{
    double start, end;
    int i, j;
    long double a[NR][NC], b[NR][NC], c[NR][NC];
    start = MPI_Wtime();
    for (i = 0; i < NR; i++)
        for (j = 0; j < NC; j++)
            a[i][j] = i + j * 1.785;
    for (i = 0; i < NR; i++)
        for (j = 0; j < NC; j++)
            b[i][j] = i + j * 0.987;
    for (i = 0; i < NR; i++)
        for (j = 0; j < NC; j++)
        {
            c[i][j] = a[i][j] * b[i][j];
        }
    printf("\nResultant Matrix:\n");
    for (i = 0; i < NR; i++)
    {
        printf("\n");
        for (j = 0; j < NC; j++)
            printf("%3.1Lf ", c[i][j]);
    }
    printf("\nFinished.\n");
}
```

```

    end = MPI_Wtime();
    printf("\nTime= %f", end - start);
    return 0;
}

```

### Parallel Code :

```

#include "mpi.h"
#include <stdio.h>
#include <stdlib.h>
#define MASTER 0
#define FROM_MASTER 1
#define FROM_WORKER 2
#define MatA_rows 5
#define MatA_cols 5
#define MatB_cols 5
// cluster /parallel code
int main(int argc, char *argv[])
{
    int numtasks, taskid, numworkers, source, dest, mtype, rows, averow, extra, offset, i,
    j,
        k, rc;
    long double a[MatA_rows][MatA_cols], b[MatA_cols][MatB_cols],
        c[MatA_rows][MatB_cols];
    MPI_Status status;
    double start, end;
    MPI_Init(&argc, &argv);
    start = MPI_Wtime();
    MPI_Comm_rank(MPI_COMM_WORLD, &taskid);
    MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
    if (numtasks < 2)
    {
        printf("Need at least two MPI tasks\n");
        MPI_Abort(MPI_COMM_WORLD, rc);
        exit(1);
    }
    numworkers = numtasks - 1;
    char pro_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(pro_name, &name_len);
    printf("Working in Processor %s\n", pro_name);
    if (taskid == MASTER)
    {
        for (i = 0; i < MatA_rows; i++)
            for (j = 0; j < MatA_cols; j++)
                a[i][j] = i + j;
        for (i = 0; i < MatA_cols; i++)
            for (j = 0; j < MatB_cols; j++)
            {
                b[i][j] = i * j;
                c[i][j] = 0.0;
            }
    }
}

```

```

averow = MatA_rows / numworkers;
extra = MatA_rows % numworkers;
offset = 0;
mtype = FROM_MASTER;
for (dest = 1; dest <= numworkers; dest++)
{
    rows = (dest <= extra) ? averow + 1 : averow;
    MPI_Send(&offset, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);
    MPI_Send(&rows, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);
    MPI_Send(&a[offset][0], rows * MatA_cols, MPI_LONG_DOUBLE,
dest, mtype,
MPI_COMM_WORLD);
    MPI_Send(&b, MatA_cols * MatB_cols, MPI_LONG_DOUBLE, dest,
mtype,
MPI_COMM_WORLD);
    offset = offset + rows;
}
mtype = FROM_WORKER;
for (i = 1; i <= numworkers; i++)
{
    source = i;
    MPI_Recv(&offset, 1, MPI_INT, source, mtype, MPI_COMM_WORLD,
&status);
    MPI_Recv(&rows, 1, MPI_INT, source, mtype, MPI_COMM_WORLD,
&status);
    MPI_Recv(&c[offset][0], rows * MatB_cols, MPI_LONG_DOUBLE,
source, mtype, MPI_COMM_WORLD, &status);
}
printf("Result Matrix:\n");
for (i = 0; i < MatA_rows; i++)
{
    printf("\n");
    for (j = 0; j < MatB_cols; j++)
        printf("%3.2Lf ", c[i][j]);
}
end = MPI_Wtime();
printf("\nTime= %f", end - start);
}
if (taskid > MASTER)
{
    mtype = FROM_MASTER;
    MPI_Recv(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD,
&status);
    MPI_Recv(&rows, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD,
&status);
    MPI_Recv(&a, rows * MatA_cols, MPI_LONG_DOUBLE, MASTER, mtype,
MPI_COMM_WORLD, &status);
    MPI_Recv(&b, MatA_cols * MatB_cols, MPI_LONG_DOUBLE, MASTER,
mtype,
MPI_COMM_WORLD, &status);
    for (k = 0; k < MatB_cols; k++)
        for (i = 0; i < rows; i++)
        {

```

```

        c[i][k] = 0.0;
        for (j = 0; j < MatA_cols; j++)
            c[i][k] = c[i][k] + a[i][j] * b[j][k];
    }
    mtype = FROM_WORKER;
    MPI_Send(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);
    MPI_Send(&rows, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);
    MPI_Send(&c, rows * MatB_cols, MPI_LONG_DOUBLE, MASTER, mtype,
            MPI_COMM_WORLD);
}
MPI_Finalize();
}

```

## Output :

```

mpiuser@c01: ~/mirror/Lab-4
File Edit View Search Terminal Help
mpiuser@c01:~/mirror/Lab-4$ mpicc parallel.cpp
mpiuser@c01:~/mirror/Lab-4$ for i in {2,4,6,8,12,16,20,32,64,128}; do mpirun -n $i -f machinefile ./a.out; done
For Number of tasks = 2    Time = 0.010618 seconds
For Number of tasks = 4    Time = 0.018650 seconds
For Number of tasks = 6    Time = 0.040115 seconds
For Number of tasks = 8    Time = 0.045943 seconds
For Number of tasks = 12   Time = 0.095248 seconds
For Number of tasks = 16   Time = 0.266002 seconds
For Number of tasks = 20   Time = 0.311706 seconds
For Number of tasks = 32   Time = 1.068787 seconds
For Number of tasks = 64   Time = 13.463950 seconds
For Number of tasks = 128  Time = 65.988798 seconds
mpiuser@c01:~/mirror/Lab-4$

```

## Compilation and Execution:

Compiling using mpic++ parallel.cpp

For execution use

```
for i in {2,4,6,8,12,16,20,32,64,128}; do mpirun -n $i -f machinefile ./a.out
```

**Observations:**

Number of Threads	Execution Time	Speed-up	Parallelization Fraction
1	0.008	1	
2	0.010	0.8	-0.5
4	0.018	0.4444	-1.667
6	0.040	0.2	-4.8
8	0.045	0.1778	-5.2849
12	0.095	0.0842	-11.8653
16	0.266	0.0301	-34.3708
20	0.311	0.0257	-39.9058
32	1.068	0.0075	-136.6022
64	13.46	0.0006	-1692.1058
128	65.98	0.0001	-10077.7323

Speed up can be found using the following formula,

$$S(n)=T(1)/T(n)$$

where, S(n) = Speedup for thread count 'n'

T(1) = Execution Time for Thread count '1' (serial code)

T(n) = Execution Time for Thread count 'n' (serial code)


Parallelization Fraction can be found using the following formula,  $S(n)=1/((1 - p) + p/n)$

where, S(n) = Speedup for thread count 'n'


n = Number of threads

p = Parallelization fraction

### Number of Threads vs Execution Time:



### Number of Threads vs Speed Up:



**Inference:**

- Execution time is increasing with an increase in the number of threads.

Since the problem is of smaller complexity the overheads of parallelization seem to have more effects here.