

Report for HPC LAB

Name: Bhaskar R

Roll No: CED18I009

Programming Environment: OpenMP

Problem: Sum of N Numbers using reduction and critical section

Date: 26th August 2021

Hardware Configuration:

CPU NAME : Intel core i5 – 8250U @ 1.60 GHz

Number of Sockets : 1

Cores per Socket : 4

Threads per core : 2

L1 Cache size : 32KB (Per Core)

L2 Cache size : 256KB (Per Core)

L3 Cache size : 6MB (Shared)

RAM : 8 GB

Serial Code:

```
#include <stdio.h>
#include <time.h>
#include <omp.h>
#include <stdlib.h>
#define n 10000
int main()
{
    double a[n], rtime;
    float startTime, endTime, execTime;
    int i, k, omp_rank;
    double sum;
    int thread[] = {1, 2, 4, 6, 8, 12, 16, 20, 32, 64};
    int thread_arr_size = 10;
    sum = 0.0;
    omp_set_num_threads(thread[k]);
    startTime = omp_get_wtime();
    for (i = 0; i < n; i++)
    {
        omp_rank = omp_get_thread_num();
        a[i] = (float)i * 1.76;
        for (int j = 0; j < 100010; j++)
            a[i] = a[i] + 2;
        sum = sum + a[i];
    }
    endTime = omp_get_wtime();
    execTime = endTime - startTime;
    rtime = execTime;
    printf("rtime = %f\n", rtime);
    return 0;
}
```

Parallel Code:

```
#include <stdio.h>
#include<time.h>
#include <omp.h>
#include<stdlib.h>
#define n 10000
int main()
{
    double a[n],rtime[n],sumarr[n];
    float startTime, endTime,execTime;
    int i,k,omp_rank;
    double sum;
    int thread[]={1,2,4,6,8,12,16,20,32,64};
    int thread_arr_size=10;
    for(k=0;k<thread_arr_size;k++)
    {
        sum=0.0;
        omp_set_num_threads(thread[k]);
        startTime = omp_get_wtime();
        #pragma omp parallel private(i)
        {
            #pragma omp for reduction (+:sum)
            for(i=0;i<n;i++)
            {
                omp_rank = omp_get_thread_num();
                a[i] = (float)i*1.67 ;
                for(int j=0;j<100010;j++)
                    a[i]=a[i]+2;
                sum=sum+a[i];
            }
        }
        endTime = omp_get_wtime();
        execTime= endTime - startTime;
        rtime[k]=execTime;
        sumarr[k]=sum;
    }
    for (k=0;k<thread_arr_size;k++)
        printf("\nThread=%d\t rtime=%f\n",thread[k],rtime[k]);
    return 0;
}
```

Compilation and Execution:

For enabling OpenMP environment use -fopenmp flag while

compiling using g++. **g++ -fopenmp sum.cpp**

For execution use

./a.out

Observations:

Number of Threads	Execution Time	Speed-up	Parallelization Fraction
1	3.335938	1	
2	1.544922	2.15	106.9
4	1.070312	3.11	90.4
6	0.712891	4.67	94.3
8	0.589844	5.65	94.0
12	0.615234	5.42	88.9
16	0.599609	5.56	87.4
20	0.583984	5.71	86.8
32	0.539062	6.18	86.5
64	0.609375	5.47	83.0

Speed up can be found using the following formula,

$$S(n)=T(1)/T(n)$$

where, S(n) = Speedup for thread count 'n'

T(1) = Execution Time for Thread count '1' (serial code)

T(n) = Execution Time for Thread count 'n' (serial code)

Parallelization Fraction can be found using the following formula, **$S(n)=1/((1 - p) + p/n)$**

where, S(n) = Speedup for thread count 'n'

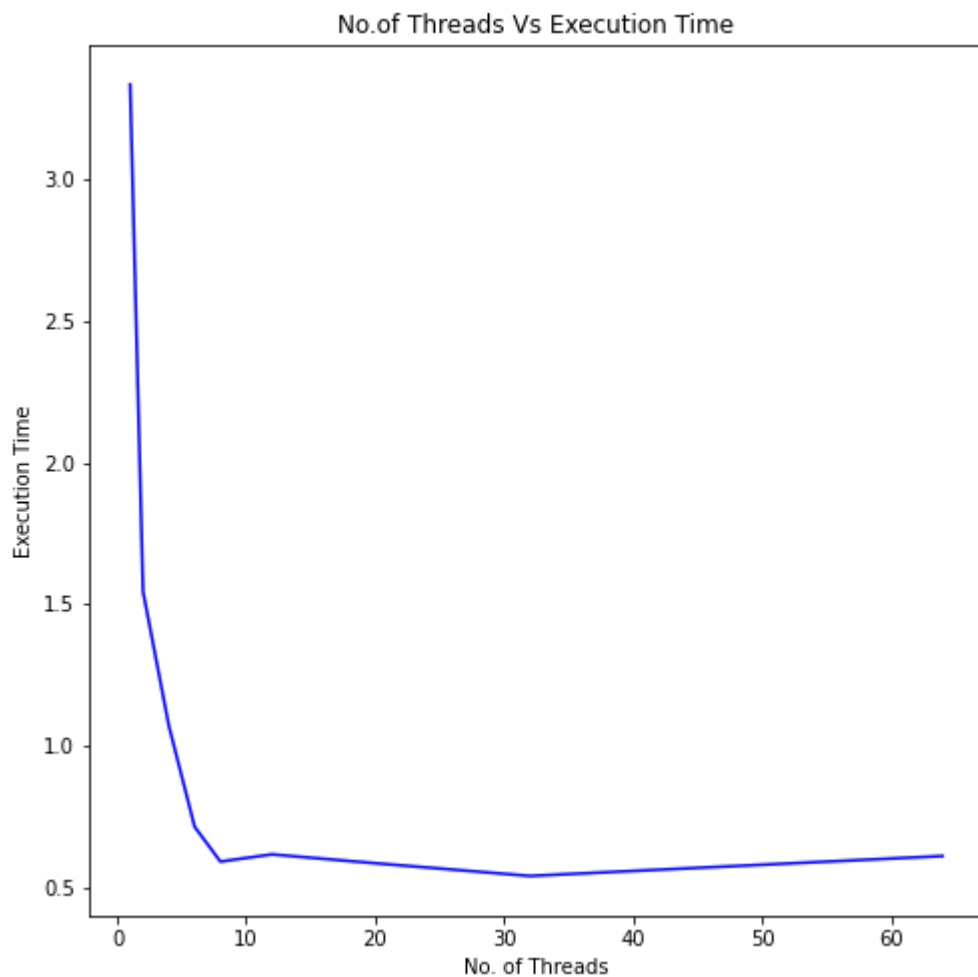
n = Number of threads

p = Parallelization fraction


Assumption:

Following extra for loop is added to increase the number of operations in the parallel region to visualize the effect of multi-threading in sum.

```
for (int j = 0; j < 100010; j++)  
    a[i] = a[i] + 2;  
sum = sum + a[i];
```

Number of Threads vs Execution Time:

Number of Threads vs Speed Up:



Inference:

(Note: Execution time, graph and inference will be based on hardware configuration)

- At thread count 32 maximum speedup is observed as the maximum number of parallel threads supported by the hardware is 8.
- If the thread count is more than 8 then the execution time increases slightly and tapers out after 12 threads.

ii) Parallel Code:(Critical Section)

```
#include <stdio.h>
#include <time.h>
#include <omp.h>
#include <stdlib.h>
#define N 100000
#define DELAY 100000
int threads[] = {1, 2, 4, 6, 8, 10, 12, 16, 20, 32, 64};
double rtime[11];
int main()
{
    double fsum, psum, a[N];
    double start, end;
    int i, k, l;
    for (k = 0; k < 11; k++)
    {
        omp_set_num_threads(threads[k]);
        fsum = 0.0;
        start = omp_get_wtime();
#pragma omp parallel
        {
            double psum = 0.0;
#pragma omp for
            for (i = 0; i < N; i++)
            {
                a[i] = (float)(i % 50) * 1.76;
                for (int kk = 0; kk < DELAY; kk++)
                    ;
                psum = psum + a[i];
            }
#pragma omp critical
                fsum = fsum + psum;
        }
        end = omp_get_wtime();
        rtime[k] = end - start;
        //printf("FSUM = ",fsum);
    }
    for (l = 0; l < 11; l++)
        printf("\nThread=%d\t rtime=%f\n", threads[l], rtime[l]);
    return 0;
}
```

Compilation and Execution:

For enabling OpenMP environment use -fopenmp flag while

compiling using **g++ -fopenmp sum_cs.cpp**

For execution use

./a.out

Observations:

Number of Threads	Execution Time	Speed-up	Parallelization Fraction
1	20.70	1	
2	10.37	1.99	99.49
4	5.06	4.09	100.7
6	4.24	4.88	95.4
8	4.13	5.01	91.4
12	4.37	4.73	86.0
16	4.55	4.54	83.1
20	3.96	5.22	85.0
32	4.06	5.09	82.9
64	3.95	5.24	82.2

Speed up can be found using the following formula,

$$S(n)=T(1)/T(n)$$

where, $S(n)$ = Speedup for thread count 'n'

$T(1)$ = Execution Time for Thread count '1' (serial code)

$T(n)$ = Execution Time for Thread count 'n' (serial code)

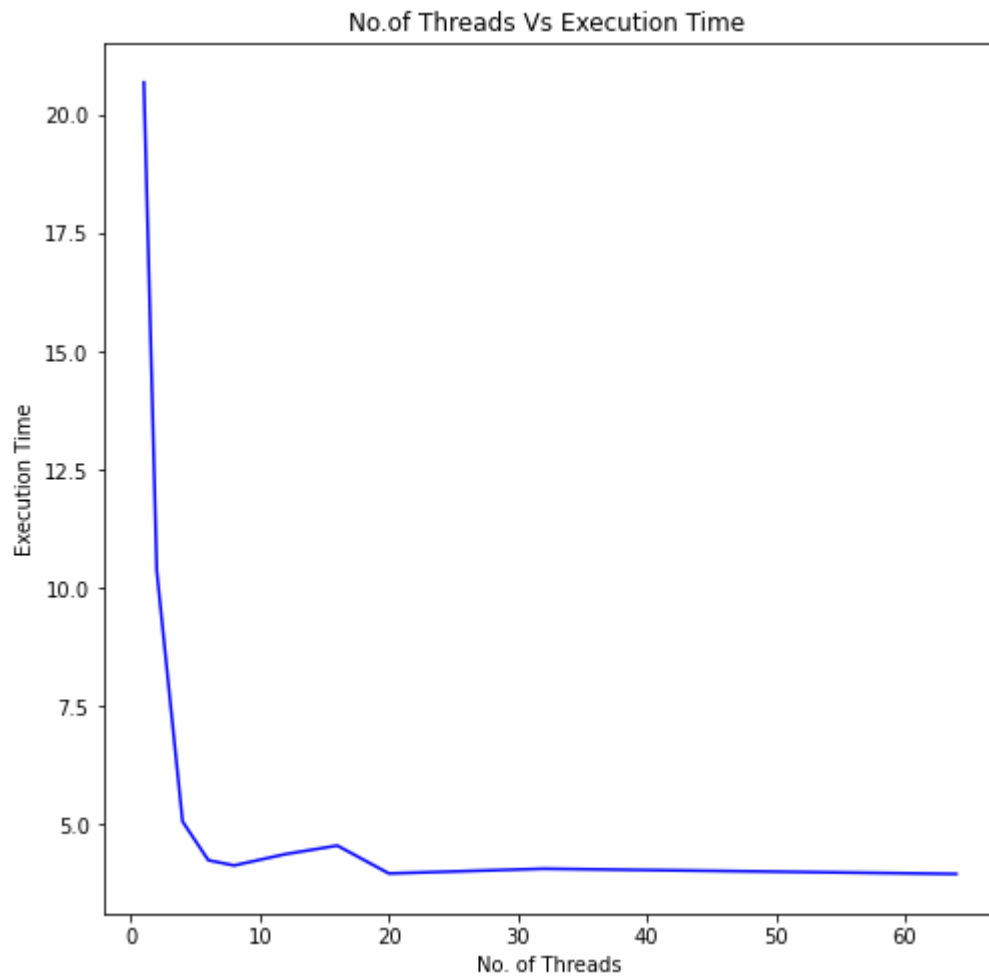
Parallelization Fraction can be found using the following formula, $S(n)=1/((1 - p) + p/n)$

where, $S(n)$ = Speedup for thread count 'n'

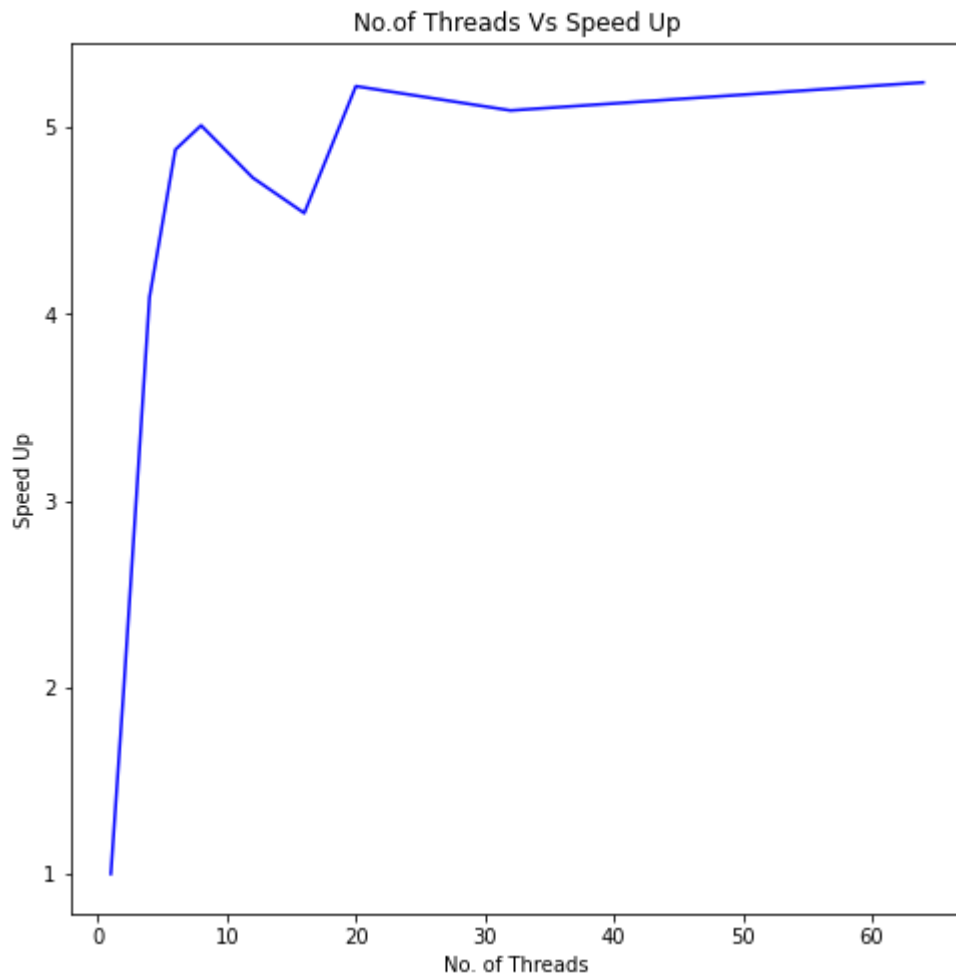
n = Number of threads

p = Parallelization fraction

Number of Threads vs Execution Time:



Number of Threads vs Speed Up:



Inference:

(Note: Execution time, graph and inference will be based on hardware configuration)

- At thread count 20 maximum speedup is observed as the maximum number of parallel threads supported by the hardware is 8.
- If the thread count is more than 10 then the execution time increases slightly and tapers out after 16 threads.