

# Report for HPC LAB

**Name:** Bhaskar R

**Roll No:** CED18I009

**Programming Environment:** OpenMP

**Problem:** Vector dot product using reduction and critical section

**Date:** 26<sup>th</sup> August 2021

## Hardware Configuration:

CPU NAME : Intel core i5 – 8250U @ 1.60 GHz

Number of Sockets : 1

Cores per Socket : 4

Threads per core : 2

L1 Cache size : 32KB (Per Core)

L2 Cache size : 256KB (Per Core)

L3 Cache size : 6MB (Shared)

RAM : 8 GB

## Serial Code:

```
#include <stdio.h>
#include <omp.h>
#define n 50000
#define delay 50000
int main()
{
    double a[n], b[n], c[n], runtime;
    float startTime, endTime;
    int i, k, omp_rank;
    double dot;
    dot = 0.0;
    startTime = omp_get_wtime();
    for (i = 0; i < n; i++)
    {
        omp_rank = omp_get_thread_num();
        a[i] = (float)i * 4.92;
        b[i] = (float)i * 2.37;
        c[i] = 0.0;
        for (int j = 0; j < delay; j++)
            c[i] += a[i] * b[i];
        dot += c[i];
    }
    endTime = omp_get_wtime();
    runtime = endTime - startTime;
    printf("runtime = %f", runtime);
    return 0;
}
```

### Parallel Code (Reduction):

```
#include <stdio.h>
#include <omp.h>

#define n 50000
#define delay 50000

int main()
{
    double a[n], b[n], c[n], runtime[11];
    float startTime, endTime;
    int i, k, omp_rank;
    double dot;
    int threads[] = {1, 2, 4, 6, 8, 10, 12, 16, 20, 32, 64};
    for (k = 0; k < 11; k++)
    {
        dot = 0.0;
        omp_set_num_threads(threads[k]);
        startTime = omp_get_wtime();
#pragma omp parallel private(i)
        {
#pragma omp for reduction(+ : dot)
            for (i = 0; i < n; i++)
            {
                omp_rank = omp_get_thread_num();
                a[i] = (float)i * 4.92;
                b[i] = (float)i * 2.37;
                c[i] = 0.0;
                for (int j = 0; j < delay; j++)
                    c[i] += a[i] * b[i];
                dot += c[i];
            }
        }
        endTime = omp_get_wtime();
        runtime[k] = endTime - startTime;
    }
    for (k = 0; k < 11; k++)
        printf("\n\nThread = %d    rtime = %f", threads[k], runtime[k]);
    return 0;
}
```

**Compilation and Execution:**

For enabling OpenMP environment use -fopenmp flag while

compiling using g++. **g++ -fopenmp dotproduct.cpp**

For execution use

./a.out

**Observations:**

Number of Threads	Execution Time	Speed-up	Parallelization Fraction
1	9.40	1	
2	4.58	2.05	102.4
4	2.90	3.24	92.1
6	3.01	3.12	81.5
8	2.62	3.58	82.3
10	2.49	3.77	81.6
12	2.60	3.61	78.8
16	3.47	2.70	67.1
20	3.45	2.72	66.5
32	3.43	2.74	65.5
64	3.50	2.68	63.6

Speed up can be found using the following formula,

$$S(n)=T(1)/T(n)$$

where, S(n) = Speedup for thread count 'n'

T(1) = Execution Time for Thread count '1' (serial code)

T(n) = Execution Time for Thread count 'n' (serial code)

Parallelization Fraction can be found using the following formula,  $S(n)=1/((1 - p) + p/n)$

where, S(n) = Speedup for thread count 'n'


n = Number of threads

p = Parallelization fraction


**Assumption:**

Following extra for loop is added to increase the number of operations in the parallel region to visualize the effect of multi-threading in the vector dot product.

```
for (int j = 0; j < delay; j++)  
    c[i] += a[i] * b[i];  
dot += c[i];
```

**Number of Threads vs Execution Time:**

### Number of Threads vs Speed Up:



### Inference:

**(Note:** Execution time, graph and inference will be based on hardware configuration)

- At thread count 10 maximum speedup is observed as the maximum number of parallel threads supported by the hardware is 8.
- If the thread count is more than 10 then the execution time increases slightly and remains the all most same after 16 threads.

---

## ii) Parallel Code:(Critical Section)

```
#include <stdio.h>
#include <omp.h>
#define n 50000
#define delay 50000
int main()
{
    double a[n], b[n], c[n], runtime[11];
    float startTime, endTime;
    int i, k, omp_rank;
    double dot, fdot;
    int threads[] = {1, 2, 4, 6, 8, 10, 12, 16, 20, 32, 64};
    for (k = 0; k < 11; k++)
    {
        dot = 0.0;
        omp_set_num_threads(threads[k]);
        startTime = omp_get_wtime();
#pragma omp parallel private(i)
        {
#pragma omp for
            for (i = 0; i < n; i++)
            {
                omp_rank = omp_get_thread_num();
                a[i] = (float)i * 4.92;
                b[i] = (float)i * 2.37;
                c[i] = 0.0;
                for (int j = 0; j < delay; j++)
                    c[i] += a[i] * b[i];
                dot += c[i];
            }
#pragma omp critical(finaldot)
                fdot += dot;
        }
        endTime = omp_get_wtime();
        runtime[k] = endTime - startTime;
    }
    for (k = 0; k < 11; k++)
        printf("\n\nThread Count: %d    Run Time: %f", threads[k], runtime[k]);
    return 0;
}
```

### Compilation and Execution:

For enabling OpenMP environment use -fopenmp flag while

compiling using **g++ -fopenmp dotproduct\_cs.cpp**

For execution use

**./a.out**

### Observations:

Number of Threads	Execution Time	Speed-up	Parallelization Fraction
1	10.53	1	
2	5.04	2.08	103
4	4.21	2.50	80
6	3.32	3.17	82.1
8	2.99	3.52	81.8
10	2.97	3.54	79.7
12	3.48	3.02	72.9
16	3.74	2.81	68.7
20	3.68	2.86	68.4
32	3.76	2.80	66.3
64	3.86	2.72	64.2

Speed up can be found using the following formula,

$$S(n)=T(1)/T(n)$$

where, S(n) = Speedup for thread count 'n'

T(1) = Execution Time for Thread count '1' (serial code)

T(n) = Execution Time for Thread count 'n' (serial code)

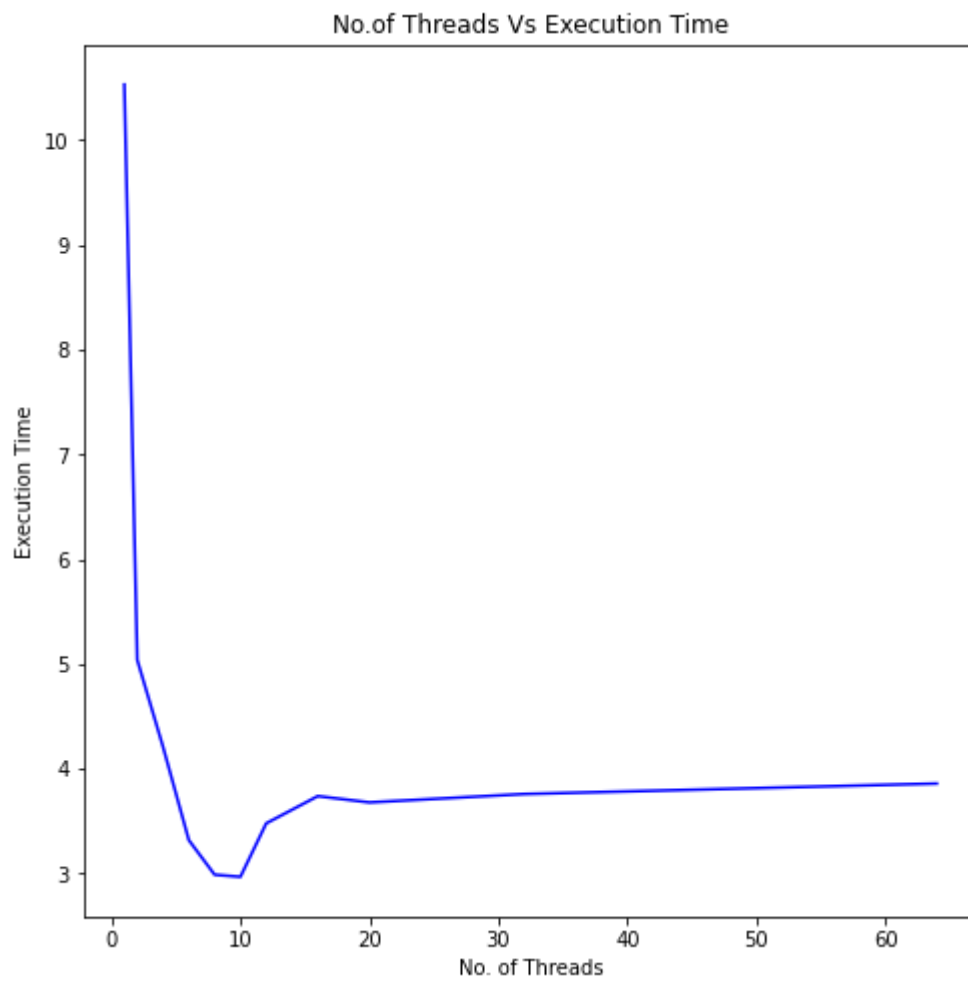
Parallelization Fraction can be found using the following formula,  **$S(n)=1/((1 - p) + p/n)$**

where, S(n) = Speedup for thread count 'n'

n = Number of threads


p = Parallelization fraction

### Number of Threads vs Execution Time:





### Number of Threads vs Speed Up:



### Inference:

**(Note:** Execution time, graph and inference will be based on hardware configuration)

- At thread count 10 maximum speedup is observed as the maximum number of parallel threads supported by the hardware is 8.
- If the thread count is more than 10 then the execution time increases slightly and remains almost the same after 16 threads.